



Deep Learning(TinyML)

Bringing Intelligence to the Edge

Chiraz Rayene Harkati
ML Engineer

December, 2025

Agenda

- AI "Born at the Edge"
- Embedded AI optimization
 - Pruning
 - Quantization
 - Distillation
 - NAS
- AI at NXP
- Lab




Image: <https://leonardocavagnis.medium.com/tinyml-machine-learning-for-embedded-system-part-i-92a34529e899>



03

AI Born at the Edge



**The ability to run
AI at the edge has
emerged as a
transformative force.**



Data privacy



Cost effective



scalable



Real-time
decisions



Energy efficiency



Human like
interactions



Contextual
understanding




Smarter,
data-driven
decisions



AI spans from training in the cloud to inference at the edge

NXP's strength in processing gives us a unique opportunity to shape the deployment of AI at the edge.



Real world examples


Cars



Waymo: 100% autonomous since 2019

Real world examples


Smartphones



FaceID
sensors



Camera image
processing



Phone
sensors



Autocorrect with language
model

Real world examples

Smart watches

- Heart rate sensor, ECG, blood oxygen...
- Applications:
 - Activity recognition
 - Sleep tracking
 - Assistant(Siri, Ok google etc)
- Everything with machine learning models!



Apple Watch and Galaxy Watch

Classical computing

A laptop or desktop CPU

- 2 to 16 cores
- Running at 2 to 5GHz
- Between 15W and 200W
- Typically, x86 instruction set
- Large chips
- Expensive (tens to hundreds of €)



Regular 15 inch HP laptop

intel **AMD**

Classical computing

A laptop or desktop CPU

- Thousands of cores
- Running at 1 to 3GHz
- Between 50W and 400W
- Extremely fast for deep learning
- Very large chips
- Very expensive (hundreds to thousands of €)




RTX 4090 Graphics card
(Nvidia)



Edge computing

ARM Microprocessors

- 1 to 8 cores
- Running at around 1GHz
- Low power (1 to 10W)
- Cheap (few €)
- Reduced instruction set (ARM)
- Typical in phones, tablets, etc.





Debix Model A with NXP i.MX 8M Plus processor

Edge computing



Microcontrollers


- 1 to 2 cores
- 10MHz to 1GHz
- Very low power (mW to 1W)
- Very cheap (few cents)
- Very reduced instruction set (sometimes no FPU)
- Typical in appliances, watches, IoT.



NXP i.MX RT Crossover
MCU

Constraints of Edge devices

- Small memory
- Small storage
- Slow(er) compute
- Limited connectivity



03

Embedded AI Optimization




Pruning





Overview



- In deep learning, pruning refers to removing unnecessary parameters that do not contribute significantly to the output.
- Motivating factors for performing pruning:
 - Identifies smaller subnetworks which can fit in memory
 - Reduces energy costs, storage constraints, inference time
 - Facilitates deployment of DL models in constrained embedded systems





What to prune

Type	Structured	Unstructured
Definition	Removing entire units (neurons, filters, channels)	Removing individual weights of neurons
Illustration		
Pros	<ul style="list-style-type: none">- Hardware acceleration- Faster inference	<ul style="list-style-type: none">- Flexibility of the method- Preservation of model structure
Cons	<ul style="list-style-type: none">- Accuracy loss- Modified structure	<ul style="list-style-type: none">- Accuracy loss- More computational cost (on normal hardware)

Global vs Local

Type	Global	Local
Definition	Removing a fixed percentage of weights from the entire model	Removing a fixed percentage of weights from each layer in the model
Illustration		
Pros	<ul style="list-style-type: none">- Easy to implement- Computationally efficient	<ul style="list-style-type: none">- Considers local weights importance- Can be adapted to the behavior of each layer
Cons	<ul style="list-style-type: none">- Important params may be removed- Loss in accuracy	<ul style="list-style-type: none">- Complex to implement and manage- Loss in accuracy




Pruning criteria

Type	Weights magnitude	Neurons activation
Definition	Removing weights with lower magnitude since they are the least important	Removing weights associated to neurons with lower activations
Illustration		
Pros	- Easy to implement	- More representative of real usage
Cons	- A layer can be pruned more than another	- Generalization problem - Requires training data

Pruning criteria

Type	Gradient based	Taylor expansion
Definition	Pruning weights with lower gradients , small updates during training	Pruning weights with lower Taylor values
Illustration		
Pros	<ul style="list-style-type: none"> - Considers direction and not only magnitude - Computationally efficient (use autograd) 	<ul style="list-style-type: none"> - Higher order information - Higher precision and accuracy
Cons	<ul style="list-style-type: none"> - Generalization problem (requires data) - Dependent to hyperparameters 	<ul style="list-style-type: none"> - Implementation complexity - Lack of scalability


When to prune?

Type	Before training	While training	After training
Definition	Pruning initialized before training the model	Pruning at specific/all training epochs	Pruning is done after the training while inference
Illustration			
Pros	<ul style="list-style-type: none">- Faster training- Smaller models	<ul style="list-style-type: none">- Adaptive, faster convergence- All computations are done before inference	<ul style="list-style-type: none">- Better preservation of accuracy- Simple and efficient
Cons	<ul style="list-style-type: none">- Risk of losing important info- Lack of prior on what to prune	<ul style="list-style-type: none">- Implementation complexity- Instability of training	<ul style="list-style-type: none">- No benefits during training- May need fine tuning

Quantization

Bit depth

- In deep learning models, parameters are represented as **floating point** (32 or 16 bit) numbers which requires a significant amount of memory and computational resources.
- Quantization **shrinks** a neural network by decreasing the **precision** of its weights, biases and activations.
- **Necessary** for some constrained devices that do not support high precision representations like the NPUs
- Pros
 - Lower latency, reduced energy
 - Computational efficiency
- Cons
 - Accuracy loss
 - Precision loss in computation



Float vs Int

0.34	3.75	5.64
1.12	2.7	-0.9
-4.7	0.68	1.43



FP32




Quantization

64	134	217
76	119	21
3	81	99

INT8




Quantization-Aware Training (QAT)



PTQ vs QAT (Source Nvidia)


PTQ vs QAT



PTQ vs QAT (Source Nvidia)

Distillation


Student Teacher



Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*.

Neural architecture search

Components of NAS



03

AI at NXP

Intelligent edge systems enabled by NXP

Expansive processor
portfolio



MCX MCUs



i.MX RT
Crossover MCUs



i.MX Apps
Processors and
beyond



AI co-processor
NPU

Differentiated HW and SW
enablement



eIQ® Neutron NPU

Highly scalable and optimized
integrated dedicated AI
acceleration

eIQ® Toolkit

AI/ML software toolkit for model
Creation, optimization and porting

eIQ® Time Series Studio

Automated ML model
creation from sensor signals

eIQ® GenAI Flow

Context aware generative AI
application development



NXP MCUs + MPUs cover the full spectrum of AI/ML applications

Legend

Rules-based /
Traditional CV

Time Series

Speech / Audio

Conv Nets (CV)

Vision
Transformers

LLMs / SLMs


VLMs / VLAs


MCX MCUs

i.MX RT Crossover MCUs

i.MX Apps Processors
and beyond


i.MX +
ARA DNPU






eIQ AI Software Toolkit

Support for MCX MCUs,
i.MX RT Crossover MCU
and i.MX Applications
Processing device
families



- NXP's **eIQ AI Software Toolkit** provides a collection of development tools, utilities and libraries for building machine learning applications using NXP MCUs and applications processors (MPUs).
- eIQ AI software can be leveraged as part of a user's existing flow or can be used for the complete flow depending on the machine learning application targeted.
- **Support** for multiple compute engines (CPU, **NPU**, GPU, DSP) if available on the NXP Processor

Addressing complexity with scalable solutions and enablement




Time Series Studio Target Applications



ANOMALY DETECTION

- Arc fault detection
- Motor health monitoring
- Water flow / leak detection
- Weight unbalance detection
- Bearing wear / gearbox malfunction
- And many more...




CLASSIFICATION

- Machine operation modes
- Presence sensing
- Appliances operation modes
- Acoustic Events
- Exercise activity types
- Gesture Types




REGRESSION

- Energy consumption prediction
- Equipment remaining useful life
- Battery Charging/Discharging prediction
- Energy demand forecasting
- Building temperature control optimization



Our GenAI solutions:


Large, Small and Vision-Language Models (LLMs, SLMs and VLMs) are creating new use-cases for edge devices



**Mobile
Robotics**




**Speech-to-speech
Interfaces**



Multi-Sense Human-Machine Interface



**Superior Reasoning
and Visual Analysis**



**Enhanced Predictive
Maintenance**

... with the introduction of new models, discrete NPUs, and multimodal AI, **the possibilities are expanding every day!**



nxp.com

| Public | NXP and the NXP logo are trademarks of NXP B.V. All other product or service names are the property of their respective owners. © 2024 NXP B.V.