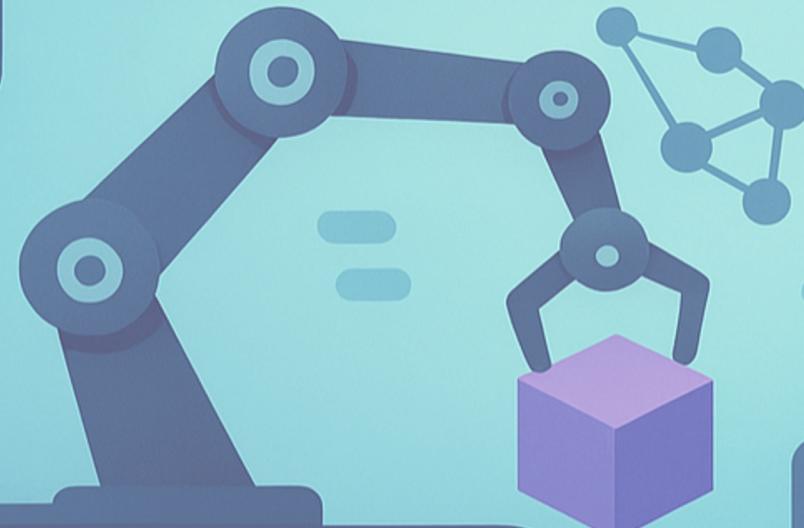


MODEL



OPTIMIZATION

APACHE TVM

References



The screenshot shows the Apache TVM Documentation homepage. The left sidebar has sections like "GETTING STARTED" (Overview, Installing TVM, Quick Start, IRModule), "HOW TO" (End-to-End Optimize Model, Customize Optimization, Optimize Large Language Model, Cross Compilation and RPC, Development Guides), "DEEP DIVE" (Design and Architecture, TensorIR, Relax, API REFERENCE, Python API, Other APIs), "ABOUT" (Contributor Guide, Publications, Security Guide), and "INDEX". The main content area is titled "Apache TVM Documentation" and includes a brief introduction, "Getting Started" (with links to Overview, Installing TVM, Quick Start, IRModule), "How To" (with links to End-to-End Optimize Model, Customize Optimization, Optimize Large Language Model, Cross Compilation and RPC, Development Guides), "Deep Dive" (with links to Overall Flow, TVM Support, TVM Runtime, TVM Node, TVM IR, TVM Target, TVM Relax, TVM TIR, TVM Arith, TVM He and TVM Topi, TVM Meta Schedule, TVM DLight, TensorIR, Tensor Program Abstraction, Understand TensorIR Abstraction, TensorIR Creation, Transformation, Relax, Graph Abstraction for ML Models, Understand Relax Abstraction, Relax Creation, Transformation), and "API Reference" (with links to Python API, Other APIs).

<https://tvm.apache.org/docs/>

The screenshot shows the microTVM documentation for TVM on bare-metal. The left sidebar includes "GETTING STARTED" (Installing TVM, Contributor Guide), "USER GUIDE" (User Tutorial, How To Guides), "DEVELOPER GUIDE" (Developer Tutorial, Developer How-To Guide), "ARCHITECTURE GUIDE" (Design and Architecture), and "TOPIC GUIDES" (microTVM: TVM on bare-metal, VTA: Versatile Tensor Accelerator, REFERENCE GUIDE, Language Reference, Python API, Other APIs, Publications, Index). The main content area is titled "microTVM: TVM on bare-metal" and discusses its purpose (running TVM models on bare-metal devices), supported hardware (STM Nucleo-F746ZG, STM32F746 Discovery, nRF 5340 Development Kit), and getting started with microTVM. It also covers "How microTVM Works" and "Help and Discussion".

<https://tvm.apache.org/docs/v0.9.0/topic/microtvm/index.html>

The slides in this presentation are highly inspired from <https://www.youtube.com/watch?v=FDmRVJkCLj4>

<https://static.linaro.org/connect/lvc21f/presentations/LVC21F-319.pdf>

https://www.edge-ai-vision.com/wp-content/uploads/2021/01/Ceze_2020_EMBEDDED_VISION_SUMMIT_Slides_Final.pdf

What is Apache TVM?



- Many models and ML frameworks landscape; models are mostly interpreted (i.e., not compiled), relying on hand-written libraries targeting specific hardware optimizations,
- Painful and unproductive for users,
- Lacks robustness for application developers whenever the model or hardware change,
- Unsustainable for hardware vendors who often develop their own software stack,
- Explosion of hardware backends (GPUs, TPU, NPUs, ASICs, FPGAs).

What is Apache TVM?

- **Apache TVM is an open-source automated end-to-end Deep Learning compiler framework for various target devices** (Apache TVM website)
 - TVM targets CPUs, GPUs, NPUs, FPGAs and μ Controllers (μ TVM),
 - Not used for training NNs,
 - Given a trained model, TVM will automatically generate and optimize tensor operators for running the model on the specified target
 - Input models in various formats, like Keras, MXNet, PyTorch, Tensorflow, Tensorflow Lite, CoreML, DarkNet, ONNX, etc.
 - Implemented in C++ and Python.

What is Apache TVM?

 TensorFlow  mxnet  PyTorch



 Keras



Automated, open source, unified optimization and compilation framework for deep learning.

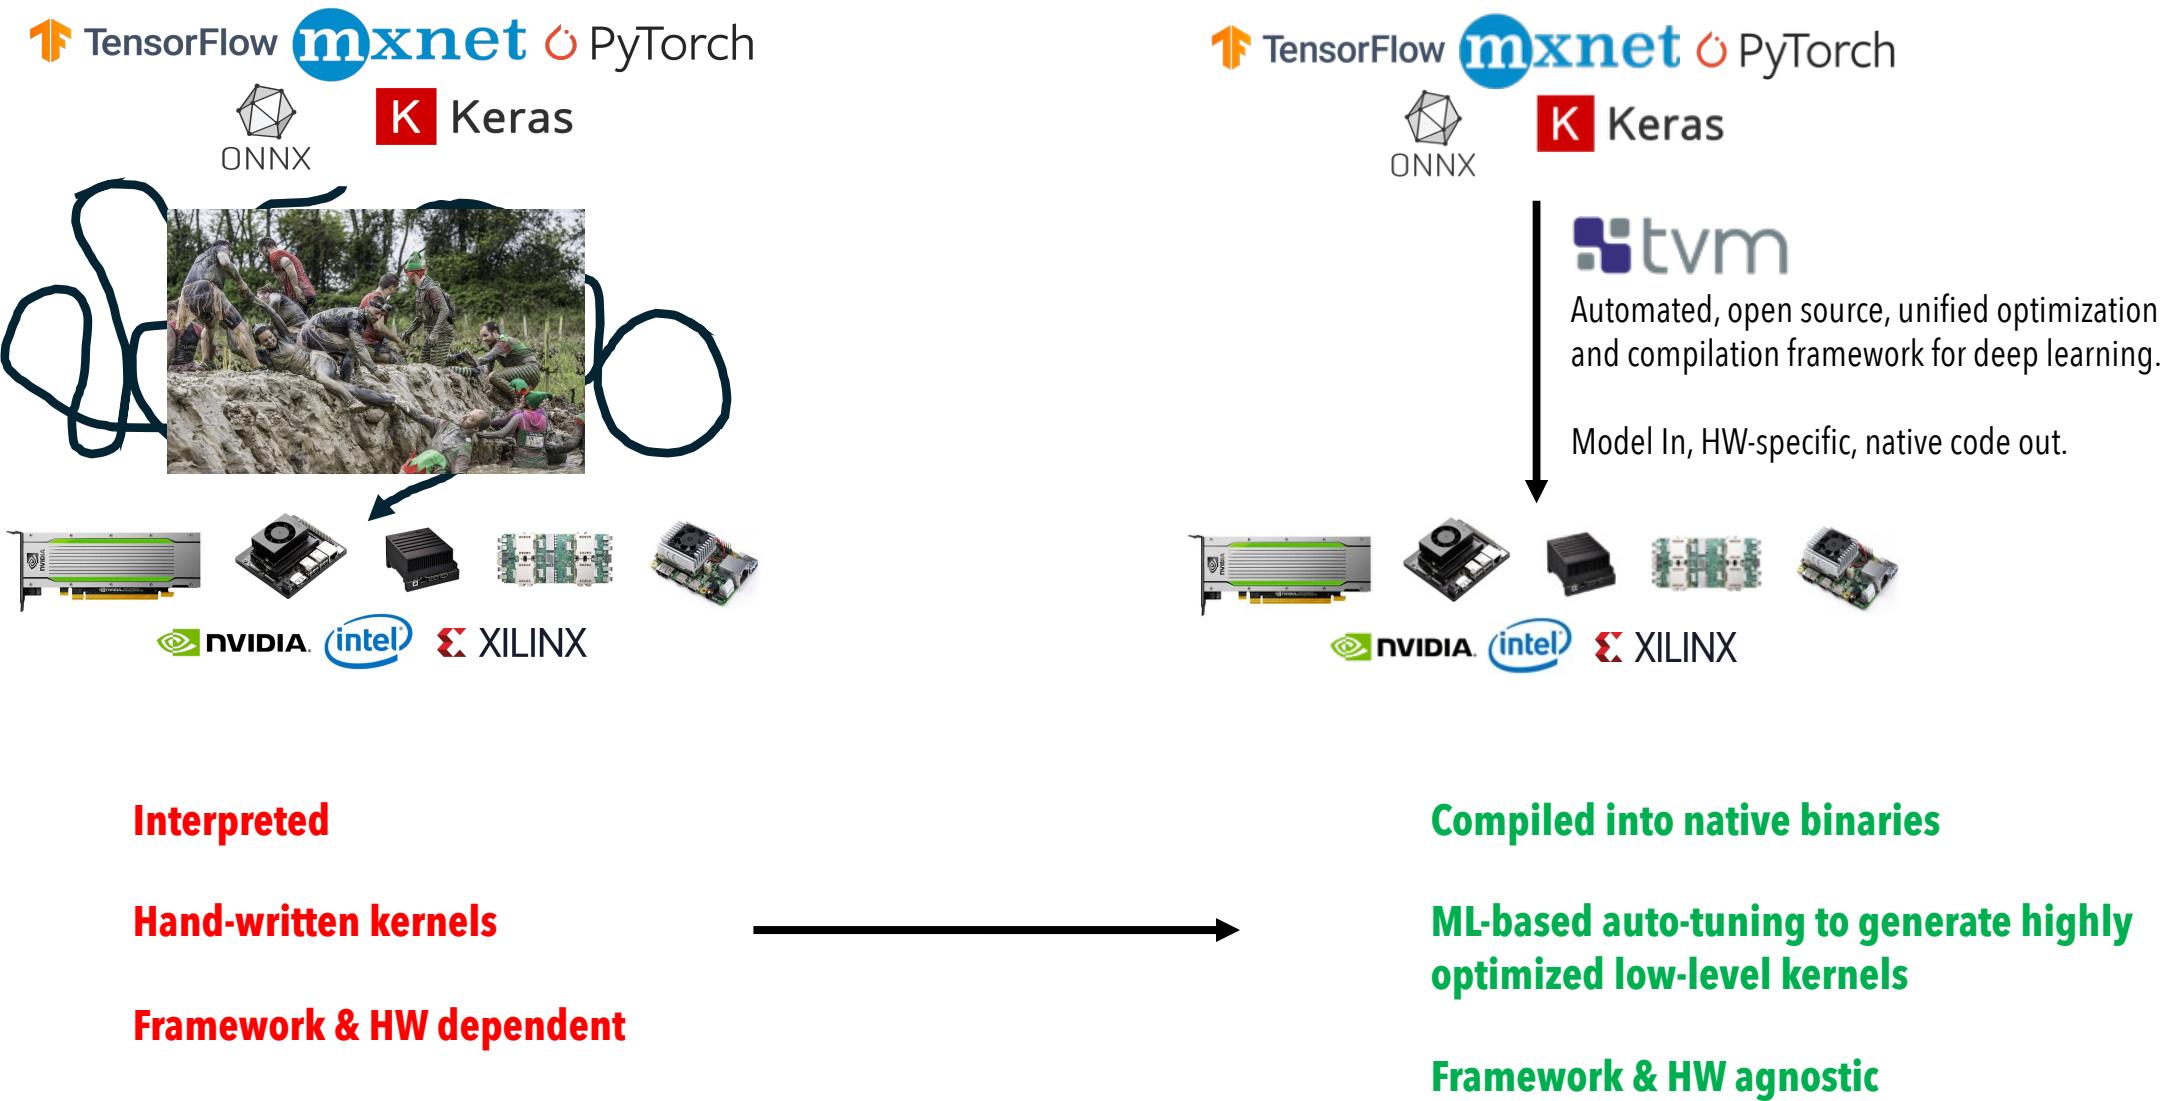
Model In, HW-specific, native code out.



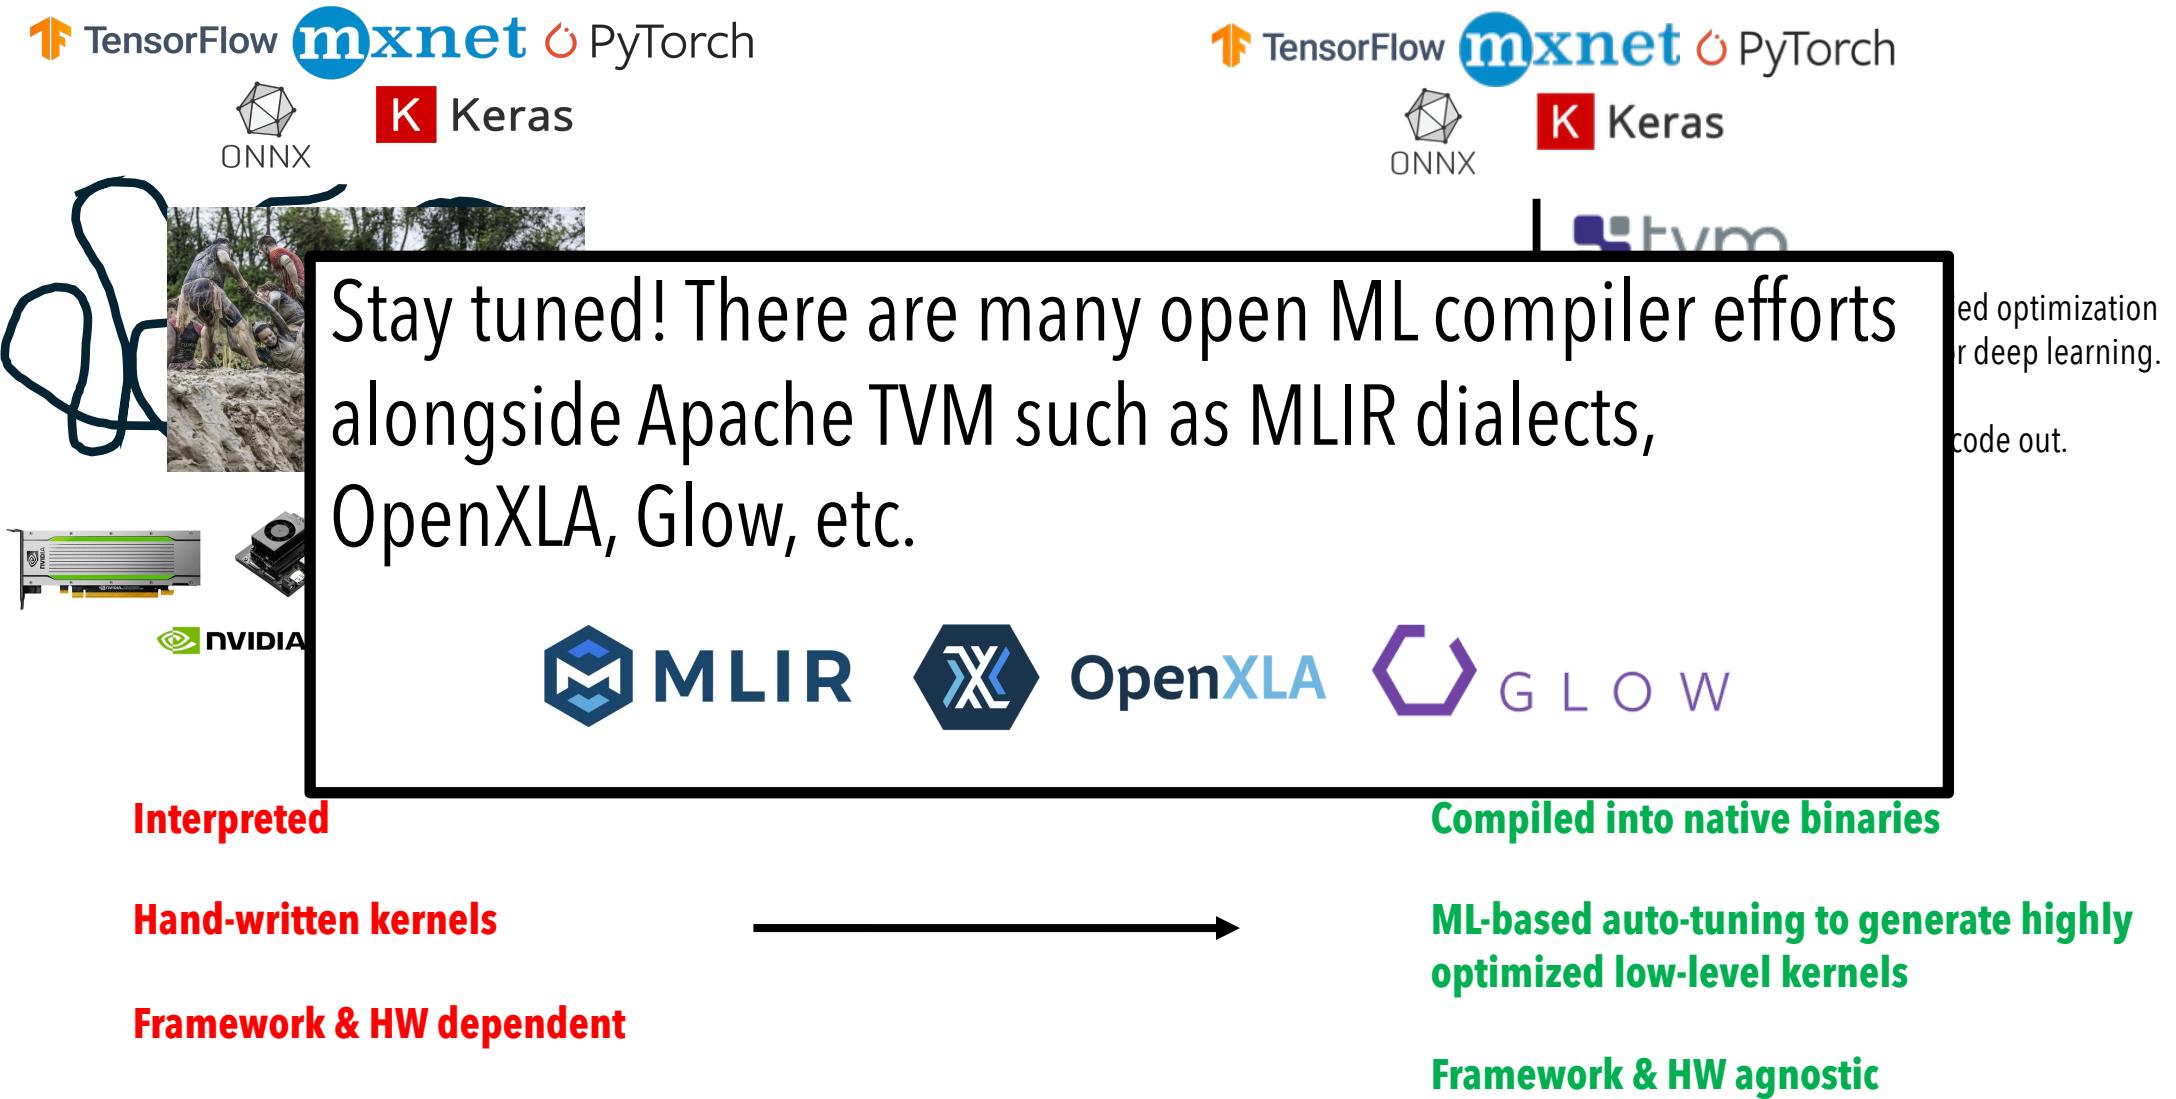
 NVIDIA  intel

 XILINX

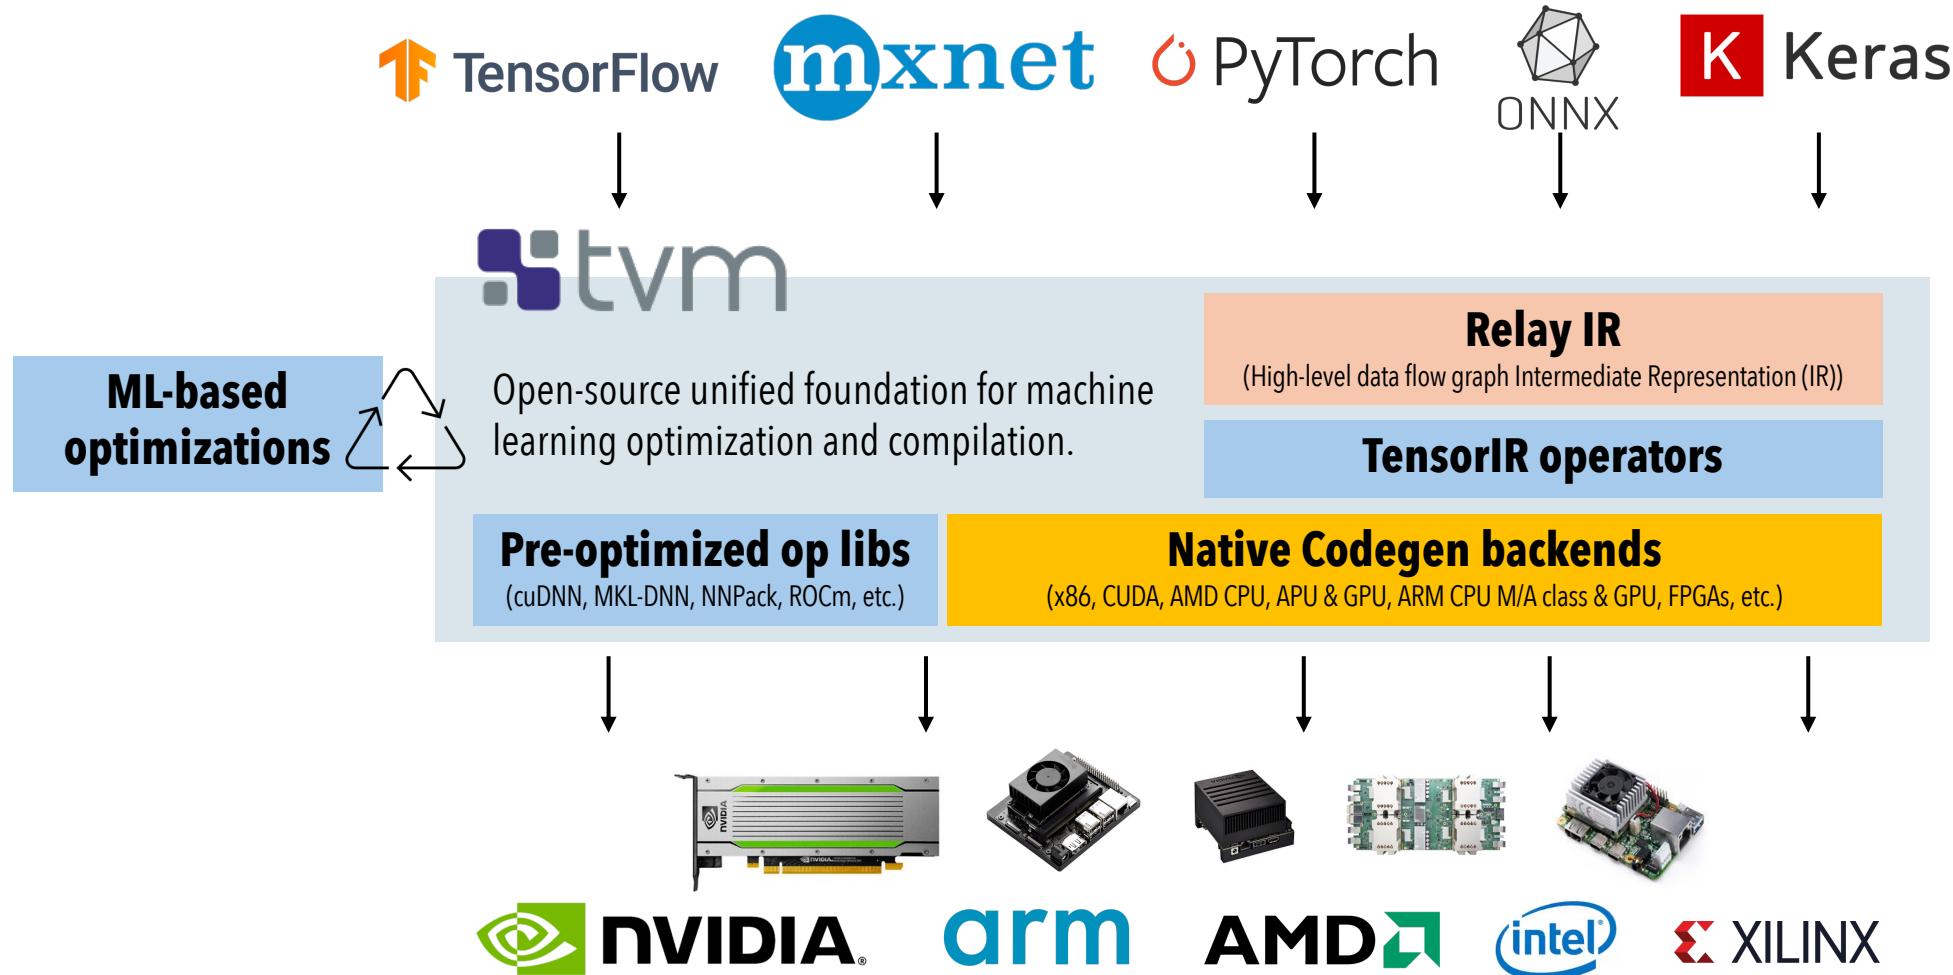
What is Apache TVM?



What is Apache TVM?



Apache TVM Stack Overview



Goals : enable short time-to-deployment for new models and new HW; easy to extend to new HW and models.

ML-based optimizations workflow

1. Model Import & Relay IR

- Model ingestion from supported frameworks,
- Conversion into Relay IR.

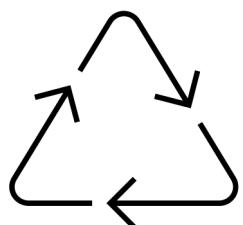
2. Operator Lowering to TensorIR

- Each operator (e.g., conv2d, matmul) is lowered to TensorIR, a loop-level intermediate representation,
- TensorIR exposes all tunable implementation choices: tiling, loop ordering, vectorization, memory layout, thread mapping, etc.
- This defines the operator-level scheduling search space.

3. ML-based Cost Model

- Instead of exhaustively benchmarking all schedules, TVM trains a cost model to predict performance of candidate schedules on the target hardware,
- Initial measurements seed the model, which is iteratively refined.

4. Auto-Tuning / MetaSchedule Loop



- Candidate TensorIR schedules are generated.
- Predicted runtime is evaluated by the cost model.
- The most promising candidates are executed on real hardware, and results feed back into the cost model.
- Loop: Generate → Predict → Measure → Update.

ML-based optimizations workflow

5. Best Schedule Selection

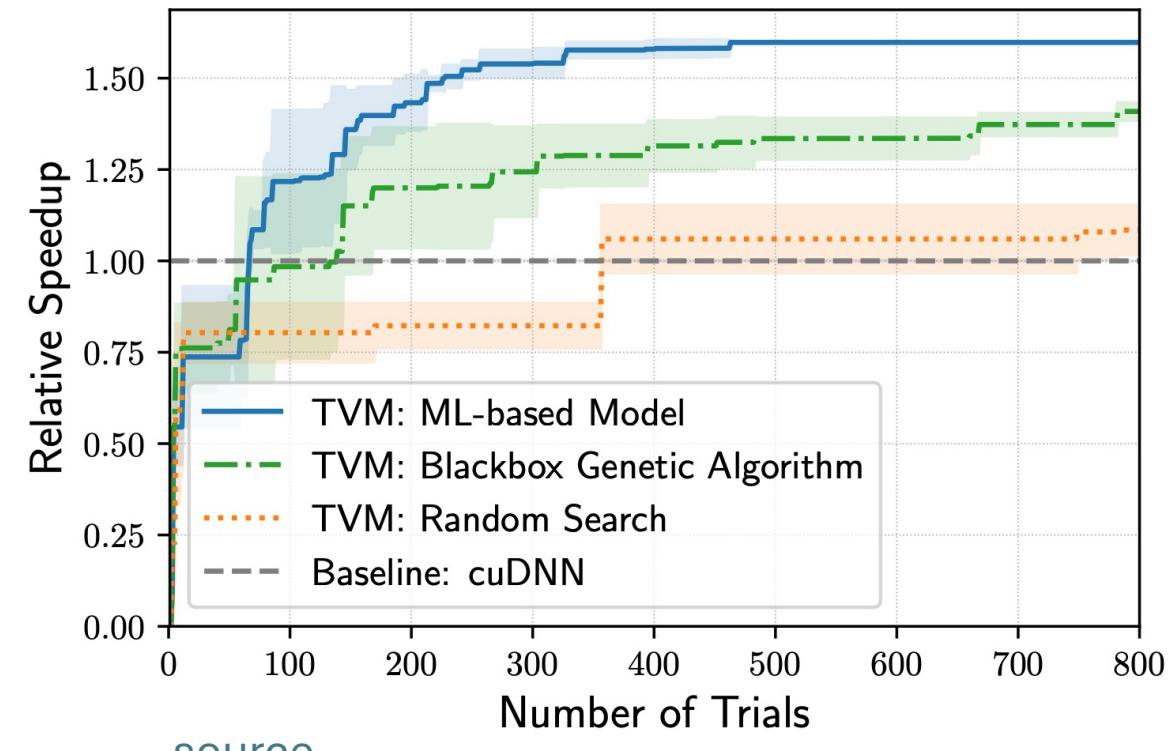
- After enough trials, the fastest schedule per operator is selected.
- Optimized kernels are cached and ready for code generation.

6. Graph-level Optimizations

- Relay-level transformations (operator fusion, quantization, memory reuse) are applied to maximize overall model performance.

7. Code Generation & Deployment

- Final optimized runtime modules are produced (CPU, GPU, or accelerator targets).
- These modules are lightweight and highly tuned for the target hardware.



[source](#)

Example (Python)

```
import tvm
from tvm import relay, auto_scheduler
from tvm.contrib import graph_executor
import onnx
import numpy as np

# Parse ONNX model
onnx_model = onnx.load("model.onnx")

input_name = "input"
input_shape = (1, 3, 224, 224)
input_dtype = "float32"

# Convert ONNX model to Relay IR
# No schedule registered for op XXX → some operators may not have an optimized kernel!
mod, params = relay.frontend.from_onnx(onnx_model, shape={input_name: input_shape}, dtype=input_dtype)

# Set hardware target and tuning options
target = tvm.target.Target("cuda") # "cuda" or "llvm" for CPU or "rocm" or "metal" or ...

# Auto-scheduler search task
tasks, task_weights = auto_scheduler.extract_tasks(mod["main"], params, target)

# Create a measure context
measure_ctx = auto_scheduler.LocalRPCMeasureContext(repeat=1, min_repeat_ms=300)

# Tuning options
tuning_option = auto_scheduler.TuningOptions(
    num_trials=100,
    measure_callbacks=[auto_scheduler.RecordToFile("tuning_records.json")],
    runner=measure_ctx.runner,
)
```

Example (Python)

```
# Run auto-tuning
for task in tasks:
    print("Tuning task:", task.name)
    tuner = auto_scheduler.TaskScheduler(task, task_weights)
    tuner.tune(tuning_option)

# Compile the model with the best schedule
with auto_scheduler.ApplyHistoryBest("tuning_records.json"):
    lib = relay.build(mod, target=target, params=params)

# Create runtime module and execute inference
dev = tvm.device(str(target), 0)
module = graph_executor.GraphModule(lib["default"])(dev)

# Run Inference
input_data = np.random.rand(*input_shape).astype(input_dtype)
module.set_input(input_name, input_data)

module.run()
output = module.get_output(0).asnumpy()
```

Example (Python)

```
# Run auto-tuning
for task in tasks:
    print("Tuning task:", task.name)
    tuner = auto_scheduler.TaskScheduler(task, task_weights)
    tuner.tune(tuning_option)
```

```
# Compile the model
with auto_scheduler.LocalRPCConfig() as config:
    lib = relay.build(mod, config=config)

# Create run function
dev = tvm.device("llvm", 0)
module = graph_runtime.create(graph, lib, dev)

with open("tuned_model_graph.json", "w") as f:
    f.write(lib.get_graph_json())

# Run Inference
input_data = np.random.uniform(0, 1, (1, 224, 224)).astype("float32")
module.set_input("input", input_data)
with open("tuned_model_params.params", "wb") as f:
    f.write(relay.save_param_dict(lib.get_params()))
module.run()
output = module["output"]()
```

One can export the compiled model

Example (CLI)

```
# Run auto-tuning
```

```
tvmc tune \  
  --target "llvm" \  
  --output tuning_records.json \  
  model.onnx
```

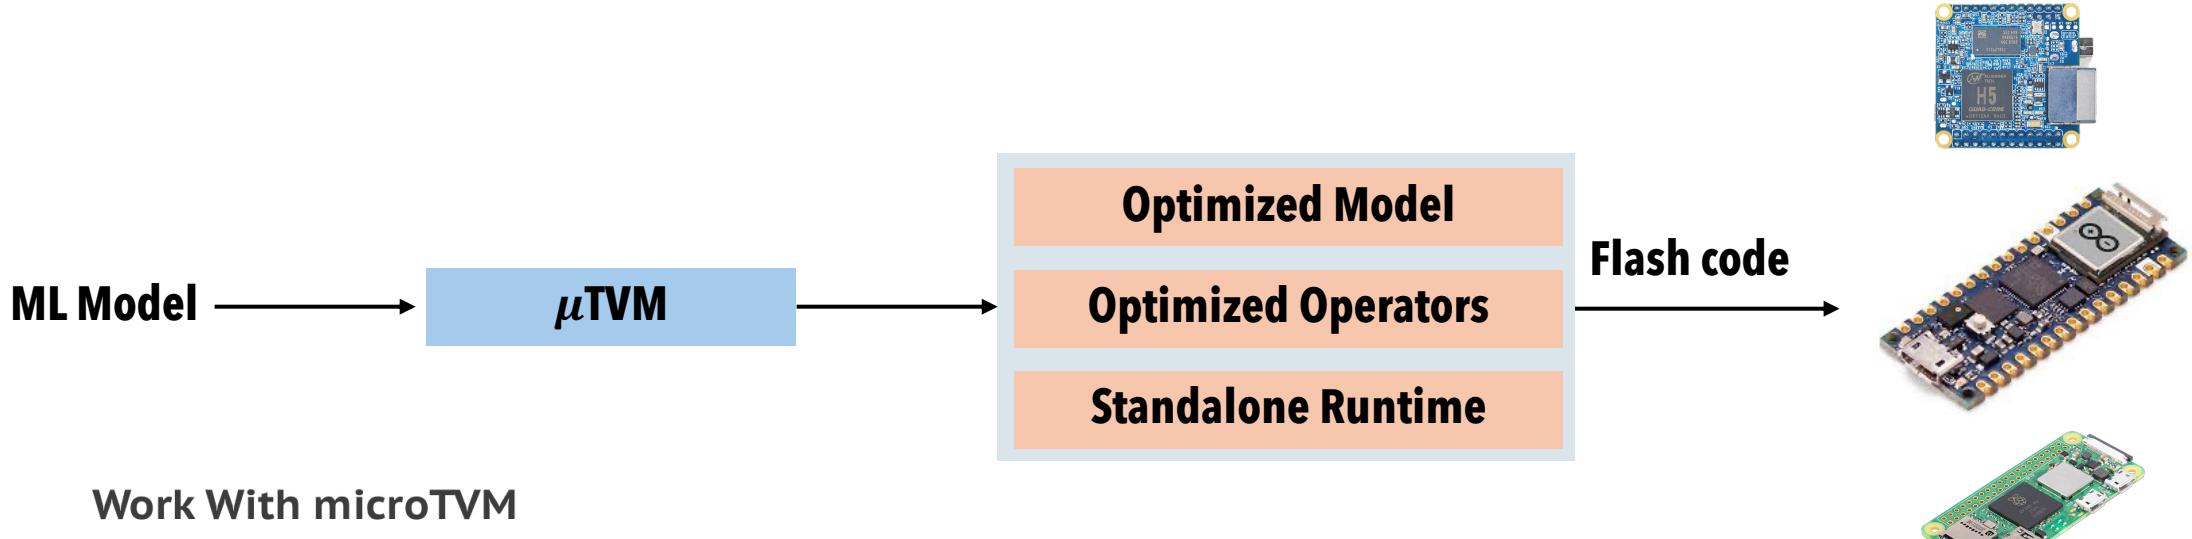
```
# Compile the model
```

```
tvmc compile \  
  --target "llvm" \  
  --tuning-records tuning_records.json \  
  --output model.tar \  
  model.onnx
```

```
# Inference
```

```
tvmc run \  
  --inputs input.npz \  
  --output predictions.npz \  
  --device cpu \  
  --print-time \  
  --repeat 100 \  
  model.tar
```

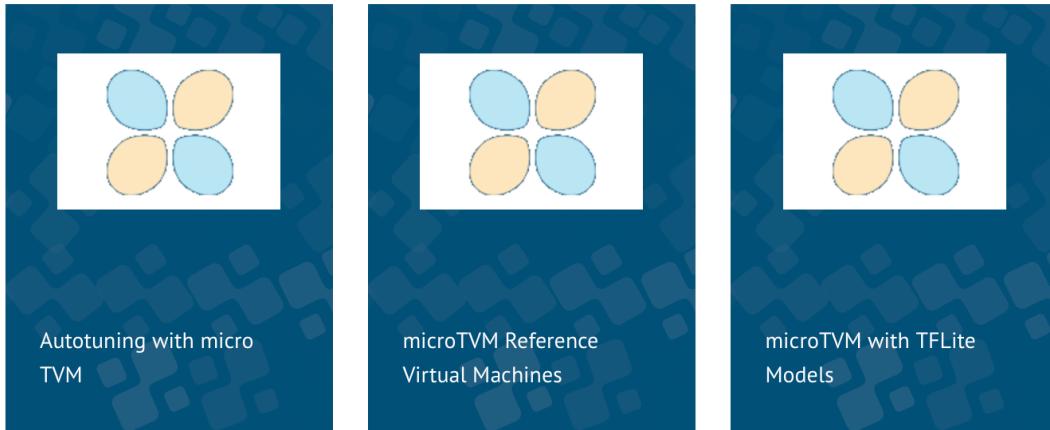
μ TVM - TVM on Bare-metal (tinyML)



Work With microTVM

microTVM enables inference on bare-metal platforms, for example, those without a traditional Operating System such as Linux, OS X, or Windows. These how-tos demonstrate how to tune and deploy models with microTVM.

Edge Devices



https://tvm.apache.org/docs/v0.8.0/how_to/work_with_microtvm/index.html

Apache TVM is an industry standard ML stack



Mobile speech recognition (85x speed-up!)



Unified ML stack for CPU, GPU, NPU built on TVM



Bing query : 112ms (TensorFlow) → 34ms (TVM)



Alexa uses a model optimized with TVM



...

Platforms Leveraging Apache TVM

NVIDIA AI Enterprise

- Integrates OctoAI, the former start-up founded by TVM developers.

truefoundry

- A cloud-native ML training and deployment platform built on Kubernetes.
- Facilitates rapid model deployment with scalability and reliability, integrating TVM for optimizations.

Kenning

- An ML framework designed for testing and deploying deep learning applications on edge devices.
- Simplifies the deployment process on edge hardware, utilizing TVM for model optimizations.

EDGE IMPULSE

- An MLOps platform tailored for developing embedded and edge machine learning systems.
- Offers tools for model optimization and deployment on a wide range of hardware targets, incorporating TVM for efficient model execution.